

INFO3406

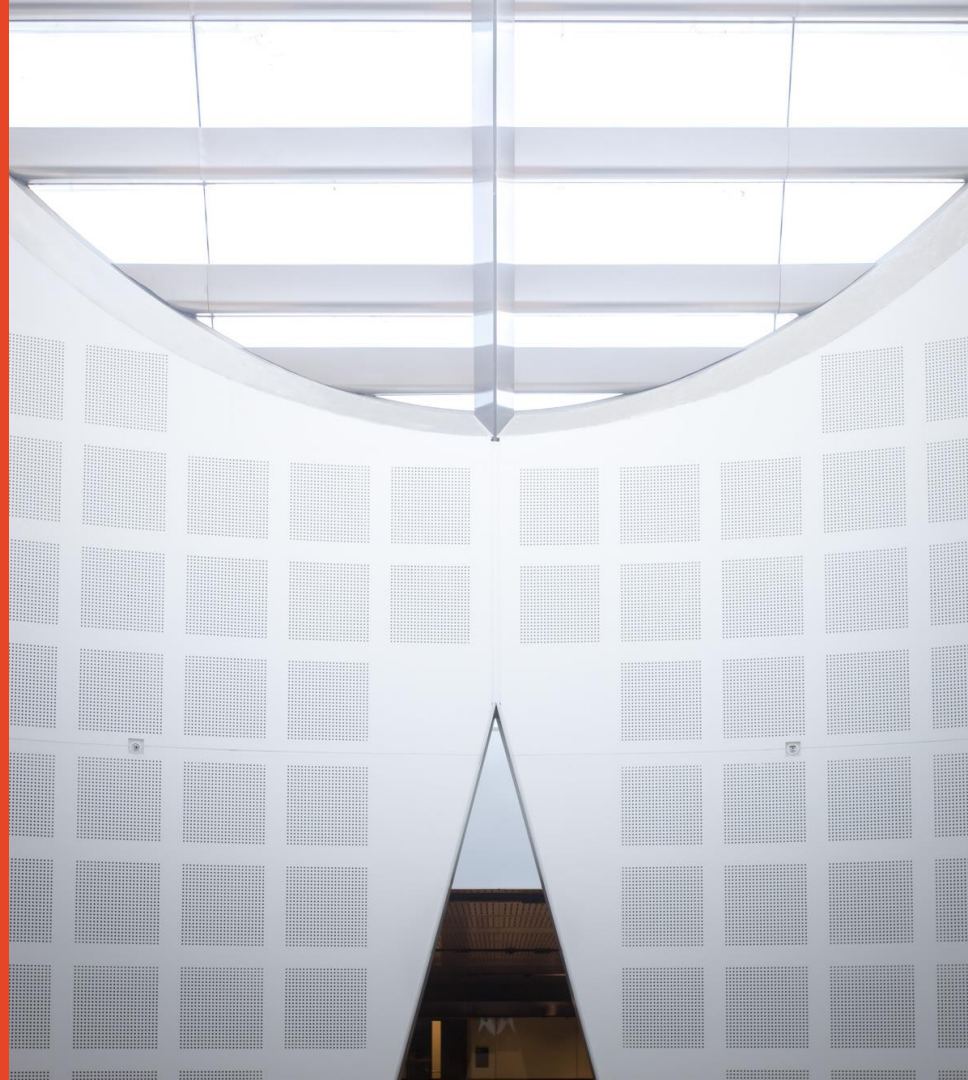
Introduction to Data Analytics

W4: Data Extract, Data Transformation and Storage

Presented by

Dr Ali Anaissi

School of IT



Overview of Week 4

Today: Data Transformation and Storage with Python and SQL

Objective

Use Python and PostgreSQL to extract, clean, transform and store data.

Lecture

- DB Access from Python
- Data cleaning and preprocessing
- Data Modeling and DB Creation
- Data Loading/Storage

Readings

- [Data Science from Scratch](#): Ch 9 + 10

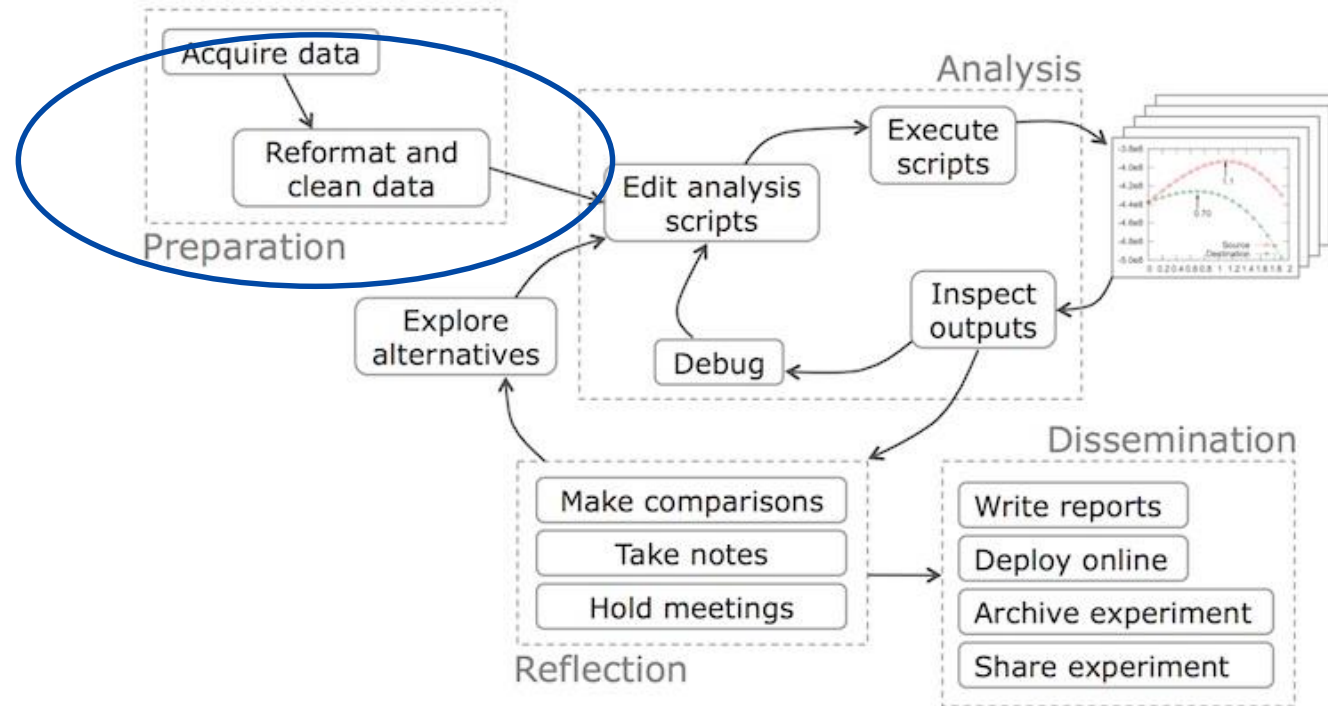
Exercises

- Python / Jupyter to load data
- psycopg2
- PostgreSQL to store data

TODO in W4

- Grok Python modules
- Grok SQL modules
- Summarise and prepare data

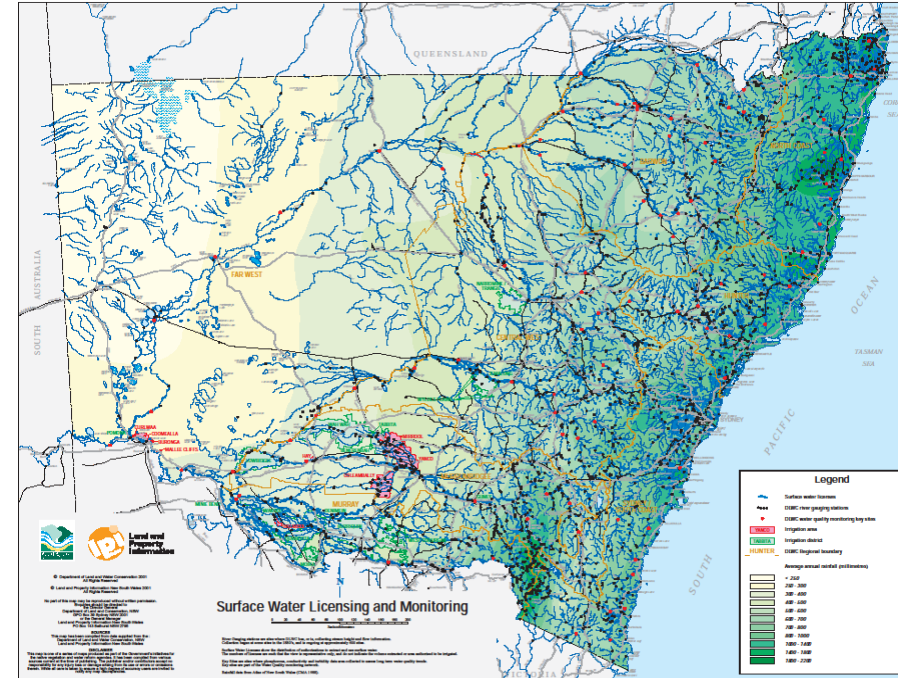
Exploratory Analysis Workflow



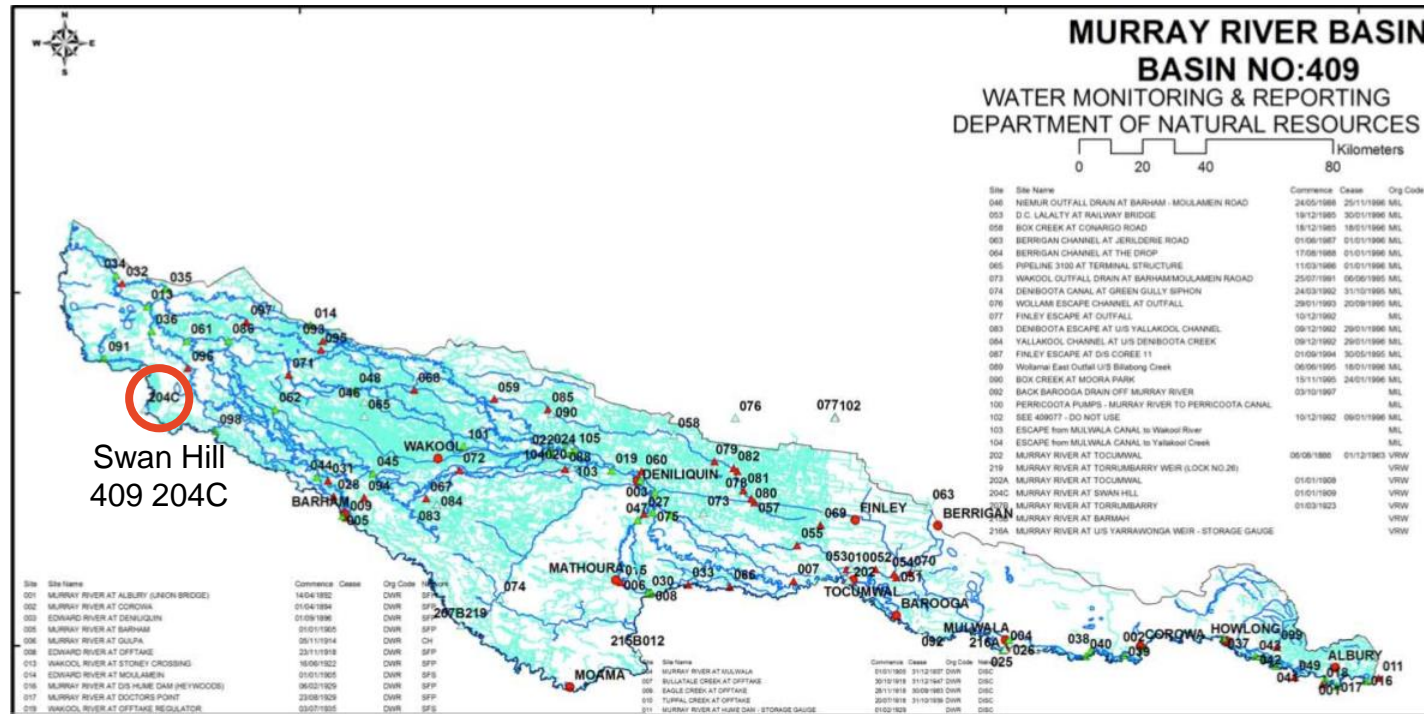
New Scenario

New Data Set

- Water measurements:
 - automatic monitoring stations that are distributed over a larger area
 - Periodically send their measured values to a central authority
 - Time-series data of:
 - water level
 - water flow
 - water temperature
 - salinity (via measuring electric conductivity) or other hydraulic properties



Example: Murray River Basin in NSW



Where do we get data from?

- You or your organization might have it already, or a colleagues provides you access to data.
 - Typical exchange formats: CSV, Excel, XML/JSON
- Or: Download from an online data server
 - Still typically in CSV or Excel etc, but now problems with meta-data
- Or: Scrap the web yourself or use APIs of resources
 - Cf. textbook, chapter 9
- Our data set comes from a colleague in Excel format

Water dataset

Contains five CSV data files:

- Measurements.csv
- Organisations.csv
- Sensors.csv
- Stations.csv

– Lets have a look

Relational Databases

- Today's goal is to store the data in a relational database
- **Relational data model** is the most widely used model today
 - Main concept: **relation**, basically a table with rows and columns
 - Every relation has a **schema**, which describes the columns, or fields
- This sounds like a spreadsheet, but as we will see, it has some differences

Definition of Relation


- **Informal Definition:**

A **relation** is a named, two-dimensional table of data

- Table consists of rows (record) and columns (attribute or field)

- Example:

Attributes (also: columns, fields)



<i>Student</i>				
<u>sid</u>	name	login	gender	address
5312666	Jones	ajon1121@cs	m	123 Main St
5366668	Smith	smith@mail	m	45 George
5309650	Jin	ojin4536@it	f	19 City Rd

Tuples
(rows,
records)



Some Remarks

- Not all tables qualify as a relation:
 - Every relation must have a unique name.
 - Attributes (columns) in tables must have unique names.
 - => The order of the columns is irrelevant.
 - All tuples in a relation have the same structure;
constructed from the same set of attributes
 - Every attribute value is atomic (not multivalued, not composite).
 - Every row is unique
(can't have two rows with exactly the same values for all their fields)
 - The order of the rows is immaterial

First attempt for database storage

- Our first attempt is for just a 1:1 mapping
 - Still: separate tables need to go to separate database relations
 - Note: CSV headers are not allowed to contain spaces or '
 - Connection to Postgresql with **psql** shell tool
 - Create 1:1 mapping tables in SQL
 - Load CSV directly to SQL tables
 - We first try COPY command from psql
- <http://www.postgresql.org/docs/current/interactive/sql-copy.html>

DB Creation and Data Import using pgsql

- Here we can work with a postgresql database using the 'psql' command.

`psql -h soit-db-pro-2.ucc.usyd.edu.au -U y18s2i3406_< your_unikey >`

- We give the following SQL create table statement:

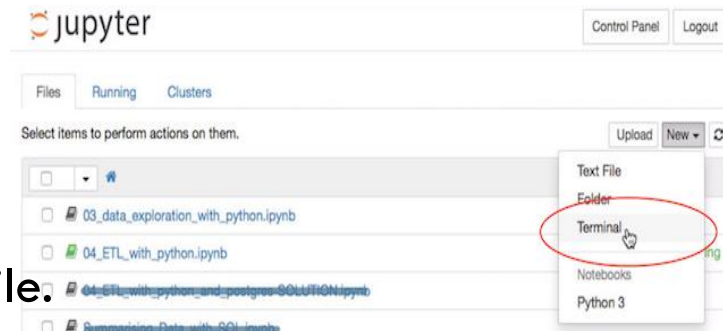
```
DROP TABLE IF EXISTS Organisation;  
CREATE TABLE IF NOT EXISTS Organisation (  
    code VARCHAR(20) PRIMARY KEY,  
    organisation VARCHAR(150));
```

- Next we want to load data from an external CSV file.

- We will use psql's `\copy` command for this.

`\copy Organisation (code,organisation) FROM 'Organisations.csv' WITH CSV HEADER`

- Psql's `\copy` command is quite useful -- as long as table and CSV files directly match, and as long as the CSV file's content is in good shape. Otherwise it soon reaches its limits.



Issues with DB Loaders

- DB Loading tools such as **psql**
 - good for administration of server
 - good for bulk-loading of exported data in clean csv format (db export)
 - needs terminal access (both an advantage and a disadvantage)
 - has its limitations if format and structure of data files do not match the actual database schema
 - => cleaning and transformation of data needed
- Could we do so in a programming language such as Python?

Database Loading with Python

Accessing PostgreSQL from Python: psycopg2

- First, we need to import the **psycopg2** module, then connect to PostgreSQL
- Note: You need obviously to provide your own login name
 - Username and password are prepared by us as your unikey and your SID

```
import psycopg2

def pgconnect():
    # please replace <your_unikey> and <your_SID> with your own details
    YOUR_UNIKEY = '<your_unikey>'
    YOUR_PW     = '<your_SID>'
    try:
        conn = psycopg2.connect(host='soit-db-pro-2.ucc.usyd.edu.au',
                                database='y18s2i3406_'+YOUR_UNIKEY,
                                user='y18s2i3406_'+YOUR_UNIKEY,
                                password=YOUR_PW)

        print('connected')
    except Exception as e:
        print("unable to connect to the database")
        print(e)
    return conn
```

Accessing PostgreSQL from Python: psycopg2 (cont'd)

- How to execute an SQL statement on an open connection 'conn'
 - we prepared a helper function which encapsulates all the error handling:

```
def pgexec( conn, sqlcmd, args, msg ):  
    """ utility function to execute some SQL statement  
        can take optional arguments to fill in (dictionary)  
        error and transaction handling built-in """  
    retval = False  
    with conn:  
        with conn.cursor() as cur:  
            try:  
                if args is None:  
                    cur.execute(sqlcmd)  
                else:  
                    cur.execute(sqlcmd, args)  
                print("success: " + msg)  
                retval = True  
            except Exception as e:  
                print("db error: ")  
                print(e)  
    return retval
```

Accessing PostgreSQL from Python: psycopg2 (cont'd)

– Example: Creating a table and loading some data

```
data_organisations = list(csv.DictReader(open('water_data/Organisations.csv')))  
# 1st: login to database  
conn = pgconnect()  
  
# 2nd: ensure that the schema is in place  
organisation_schema = """CREATE TABLE IF NOT EXISTS Organisation (  
                           code VARCHAR(20) PRIMARY KEY,  
                           orgName   VARCHAR(150)  
                           )"""  
pgexec (conn, organisation_schema, None, "Create Table Organisation")  
  
# 3rd: load data  
# IMPORTANT: make sure the header line of CSV is without spaces!  
insert_stmt = """INSERT INTO Organisation(code,orgName)  
                VALUES (%(Code)s, %(Organisation)s)"""  
for row in data_organisations:  
    pgexec (conn, insert_stmt, row, "row inserted")
```

Issues could be encountered

- Interpretation of data format and meta-data
- Differences in naming conventions
 - Excel headers with spaces and quotes, which both are not allowed to DBMS
- Inconsistent or missing data entries
- 'shape' of data

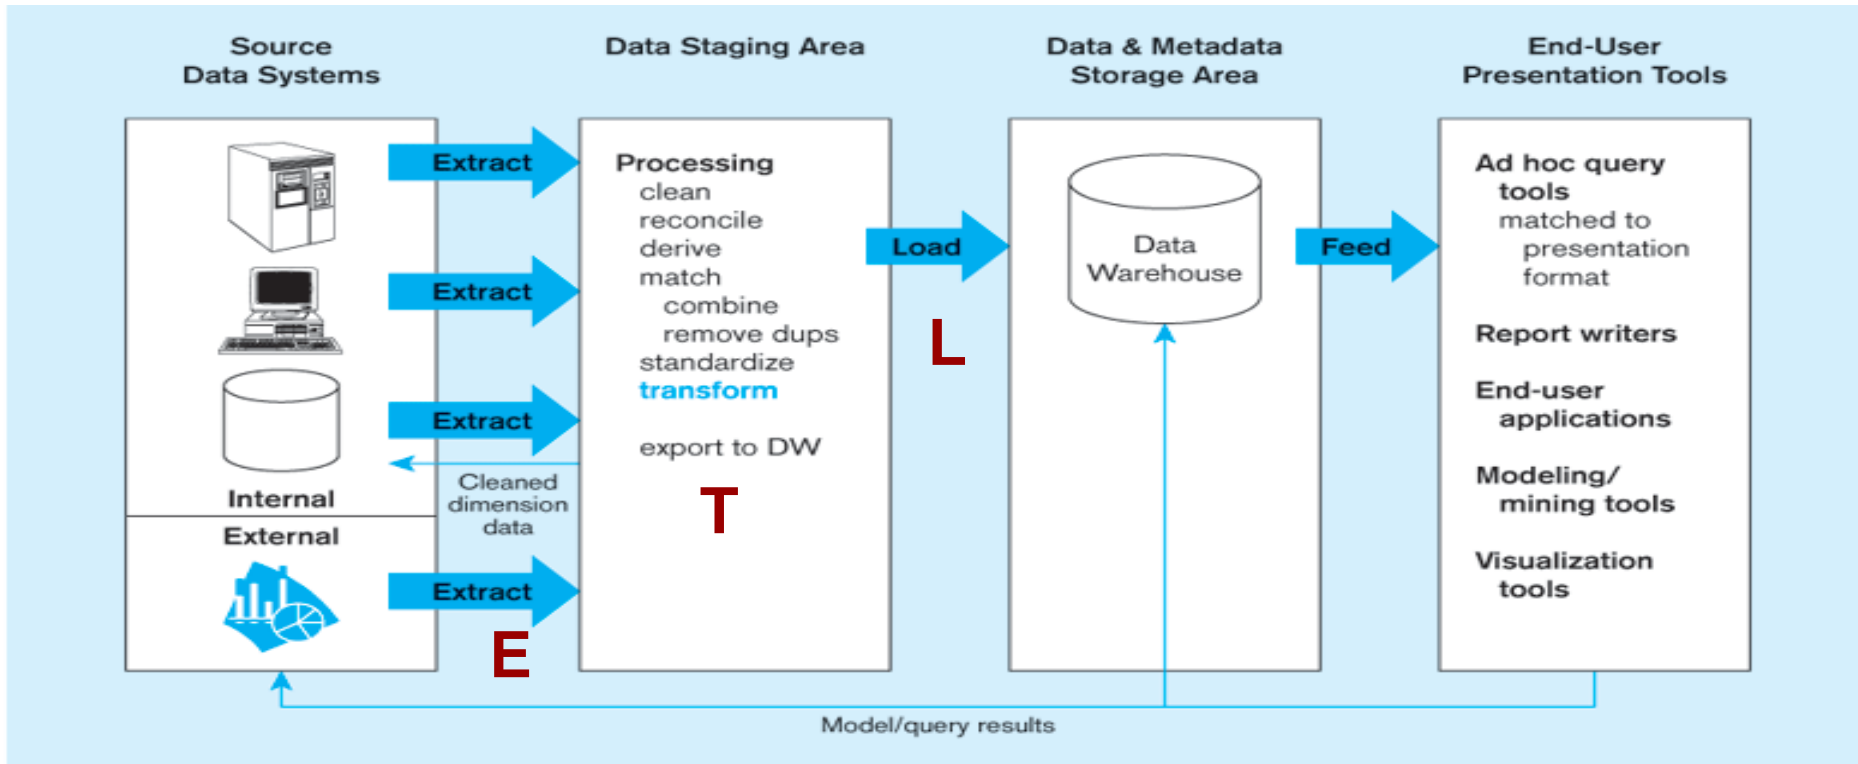
Transforming and Cleaning Data

Cleaning and Transforming Data

- Real data is often '*dirty*'
- Important to do some data cleaning and transforming first
 - remember last week, when we had to convert strings to float
- Typical steps involved:
 - type and name conversion
 - filtering of missing or inconsistent data
 - unifying semantic data representations
 - matching of entries from different sources
- Later also:
 - **Rescaling** and optional dimensionality reduction

ETL Process

- This problem is well known from data warehousing
- **ETL Process:** Capture/Extract - Data Cleansing - Transform - Load



Data Cleaning in Python

- Yes, there are powerful ETL tools out there, but we do it for free in Python:
 - (1) type and name conversion
 - (2) filtering of missing or inconsistent data
 - (3) unifying semantic data representations
 - (4) matching of entries from different sources
- Last week's `clean()` function deals with Tasks (1) and (2)
 - `int()` creates integer objects, e.g., -1, 101
 - `float()` creates floating point object, e.g., 3.14, 2.71
 - `datetime.strptime()` creates datetime objects from strings
 - Filters missing / wrongly formatted data and replaces with default value
- For more complex cases (3) and (4), you would need special code though

Repeat from Week 3: A function to convert values

Use “not a number” as default value
numpy knows to ignore for some stats

```
import numpy as np
DEFAULT_VALUE = np.nan

def clean(data, column_key, convert_function, default_value):
    special_values= {} # no special values yet
    for row in data:
        old_value = row[column_key]
        new_value = default_value
        try:
            if old_value in special_values.keys():
                new_value = special_values[old_value]
            else:
                new_value = convert_function(old_value)
        except (ValueError, TypeError):
            print('Replacing {} with {} in column {}'.format(row[column_key], new_value, column_key))
            row[column_key] = new_value

# the following converts the two measurement columns to float values - or NaN
clean(data_measurements, 'Discharge', float, DEFAULT_VALUE)
clean(data_measurements, 'MeanDischarge', float, DEFAULT_VALUE)
```

Data Modeling

Relation Database Theory and Issues

- The modelling process in relational database known as OLTP (Online Transactional Processing) focuses on normalization process which yields to a flexible model
 - making it easy to maintain dynamic relationships between business entities
- So it is effective and efficient for operational databases – a lot of updates
- However, a fully normalized data model can perform very inefficiently for queries.
- Historical data are usually large with static relationships:
 - Unnecessary joins may take unacceptably long time
- So how to proceed with a database approach?
 - => **OLAP: Online Analytical Processing (Data Warehousing Approach)**

What is a Data Warehouse?

- “A data warehouse is simply a single, complete, and consistent store of data obtained from a variety of sources and made available to end users in a way they can understand and use it in a business context.” -- Barry Devlin, *IBM Consultant*
- A data warehouse is a subject-oriented, integrated, time-variant, and nonvolatile collection of data that is used primarily in organizational decision making. -- W. H. Inmon
 - OLTP is trying to run a business, while;
 - OLAP is trying to improve/optimize a business
 - » is an element of decision support systems (DSS)

What is a Data Warehouse?

- Subject-oriented
 - Organized by subject, not by application
 - Used for analysis, data mining, etc.
- Integrated
 - Constructed by integrating multiple, heterogeneous data sources
 - relational databases, flat files, on-line transaction record
- Time Variant
 - Large volume of historical data (Gb, Tb)
 - Time attributes are important
- Non-volatile
 - Updates infrequent or does not occur
 - May be append-only

Conceptual Modeling of Data Warehouses

- Modeling data warehouses: dimensions & measures instead of relational model
- Data warehouse contains a **large** central table (**fact table**)
 - Contains the data without redundancy
- A set of **dimension tables**

Data Warehouses: Fact Tables


- Relational ‘data warehouse’ applications are centered around a ***fact table***
 - For example, a supermarket application might be based on a table
Sales (*Market_Id*, *Product_Id*, *Time_Id*, *Sales_Amt*)

market_id	product_id	time_id	sales_amt
M1	P1	T1	3000
M1	P2	T1	1000
M1	P3	T1	500
M2	P1	T1	100
M2	P2	T1	1100
M2	P3
...	...		

- The table can be viewed as *multidimensional*
 - Collection of numeric measures, which depend on a set of dimensions
 - E.g. *Market_Id*, *Product_Id*, *Time_Id* are the dimensions that represent specific supermarkets, products, and time intervals
 - *Sales_Amt* is a function of the other three

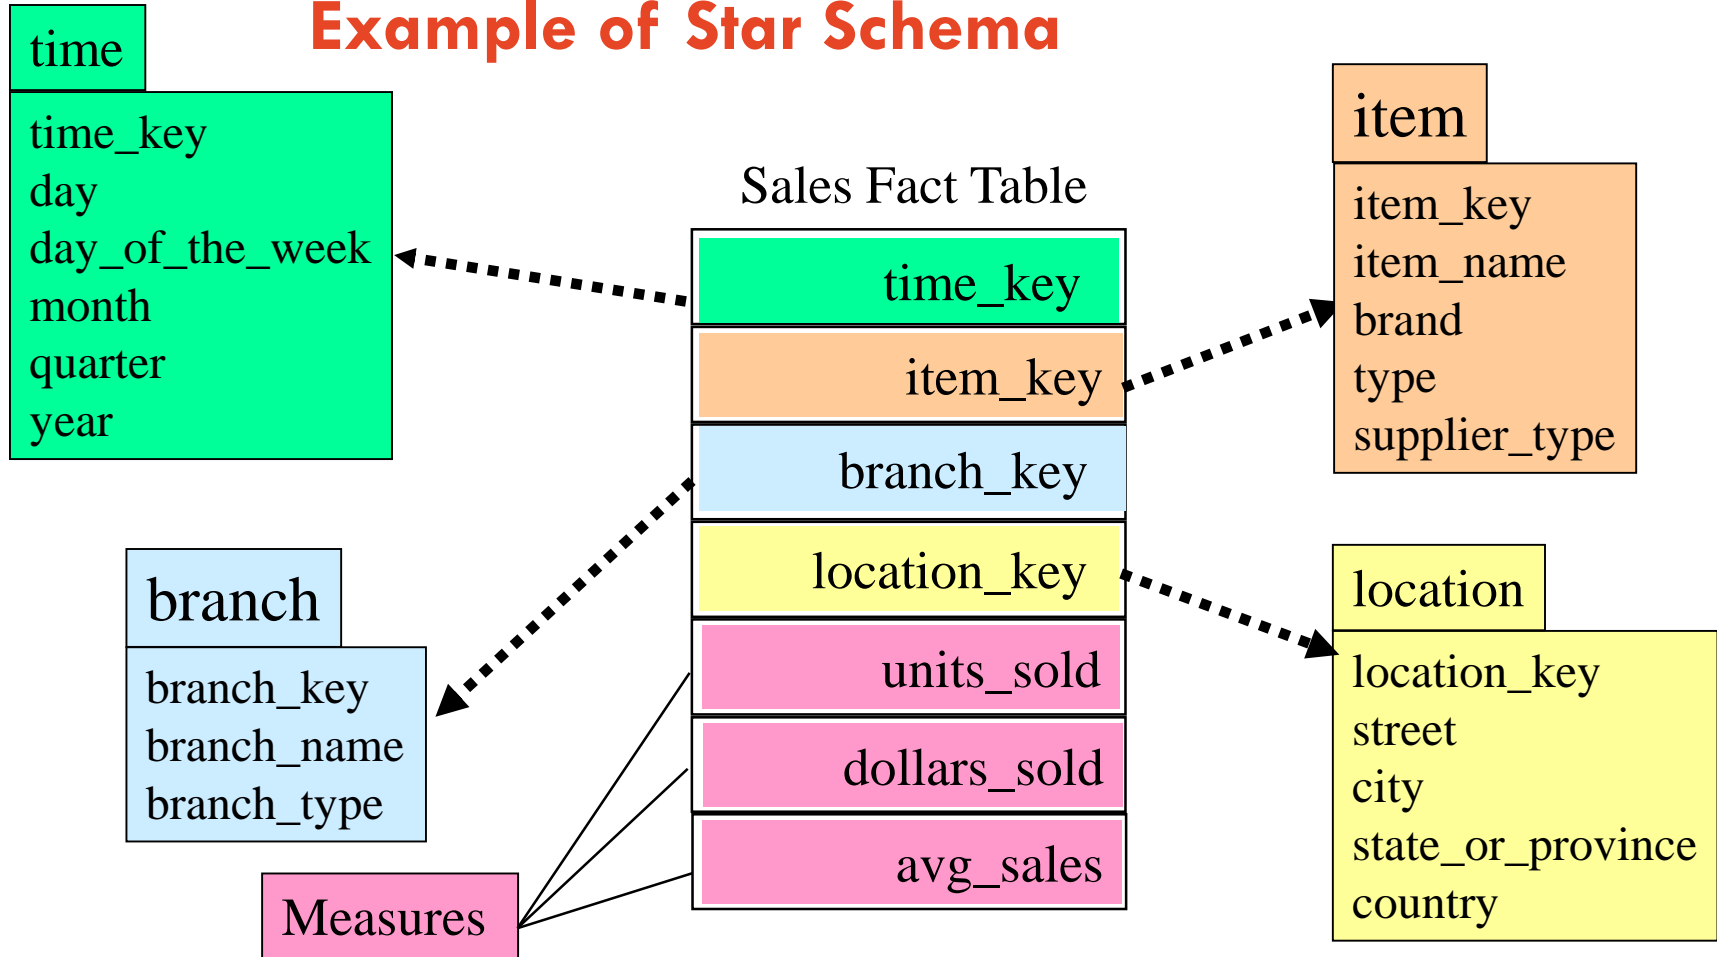
Data Warehousing: Star Schema

- The fact and dimension relations linked to it looks like a star;
- this is called a **star schema**
- *Most common modeling paradigm*




- If we map this to relations
 - 1 central fact table
 - n dimension tables with foreign key relationships from the fact table
- (the fact table holds the FKs referencing the dimension tables)*

Example of Star Schema

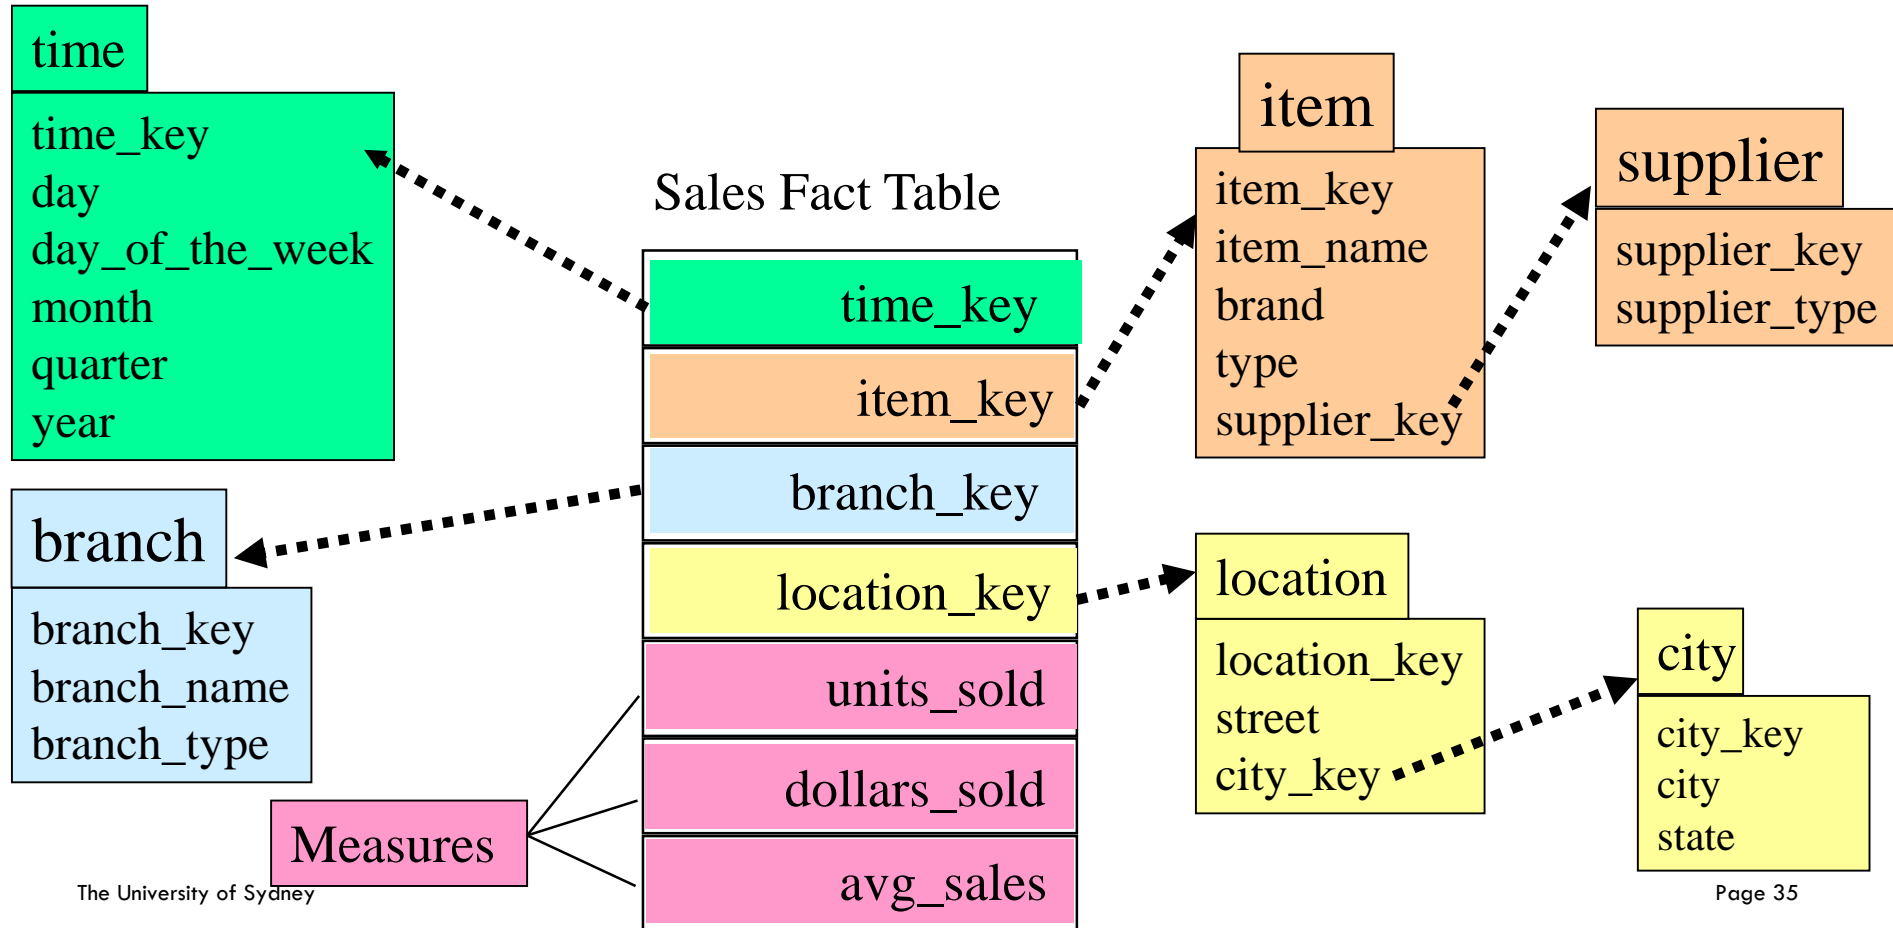


Data Warehousing: Snowflake Schema

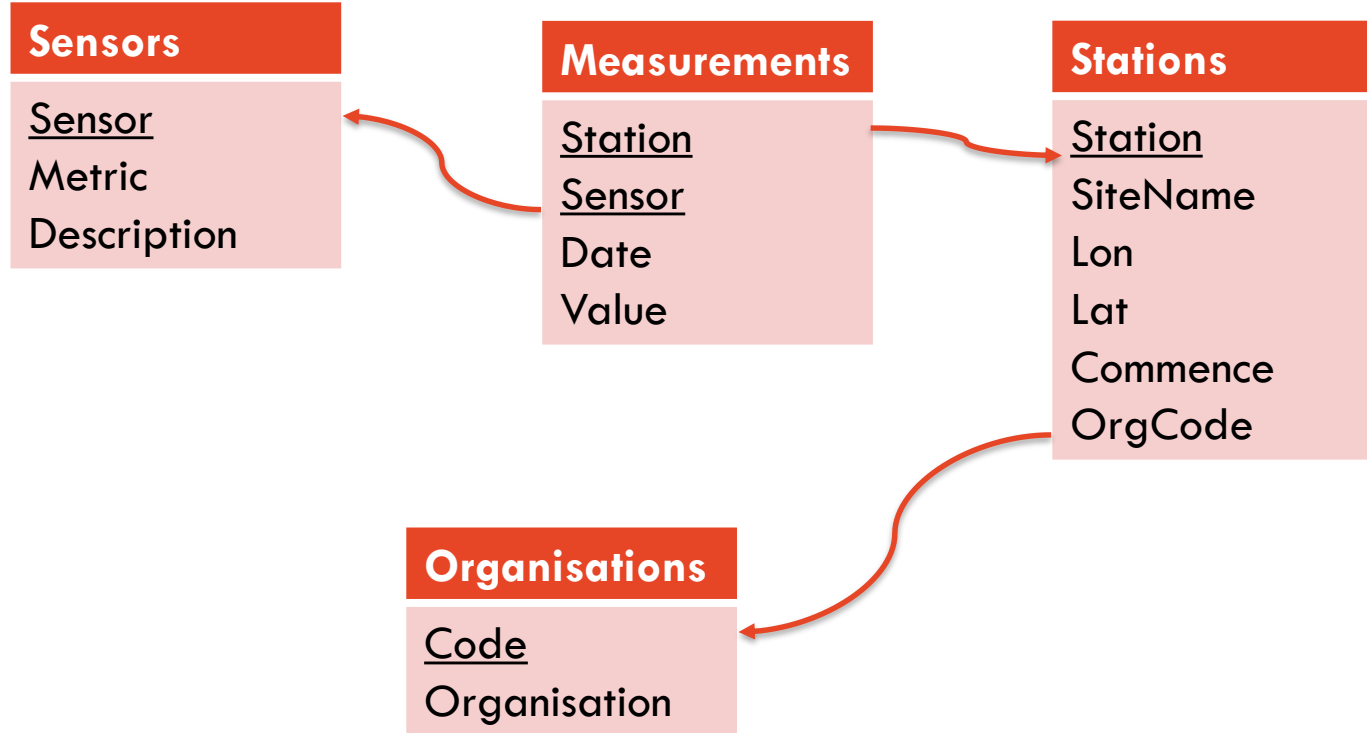
- Snowflake schema: A refinement of star schema where some dimensional hierarchy is **normalized** into a set of smaller dimension tables, forming a shape similar to snowflake
- measurements are the facts, rest describes the dimensions



Example of Snowflake Schema



Modeling our Water Data Set



Data Warehousing: Fact constellations

- Fact constellations: **Multiple** fact tables share dimension tables, viewed as a collection of stars, therefore called galaxy schema or fact constellation

DB Creation

SQL – The Structured Query Language

- SQL is the standard declarative query language for RDBMS
- Supported commands from roughly two categories:
 - **DDL** (Data Definition Language)
 - Create, drop, or alter the relation schema
 - Example:
CREATE TABLE *name* (*list_of_columns*)
 - **DML** (Data Manipulation Language)
 - for retrieval of information also called **query language**
 - **INSERT, DELETE, UPDATE**
 - **SELECT ... FROM ... WHERE**

Table Constraints and Relational Keys

- When creating a table, we can also specify **Integrity Constraints** for columns
 - eg. domain types per attribute, or **NULL** / **NOT NULL** constraints
- **Primary key:** unique, minimal identifier of a relation.
 - Examples include employee numbers, social security numbers, etc. This is how we can guarantee that all rows are unique.
- **Foreign keys** are identifiers that enable a dependent relation (on the many side of a relationship) to refer to its parent relation (on the one side of the relationship)
 - Must refer to a candidate key of the parent relation
 - Like a 'logical pointer'
- Keys can be **simple** (single attribute) or **composite** (multiple attributes)

Example: Relational Keys

Primary key identifies each tuple of a relation.

Composite Primary Key consisting of more than one attribute.

<i>Student</i>	
<u>sid</u>	name
31013	John

<i>Enroll</i>		
<u>sid</u>	<u>ucode</u>	grade
31013	I2120	CR

<i>Units_of_study</i>		
<u>ucode</u>	title	credit_pts
I2120	DB Intro	4

Foreign key is a (set of) attribute(s) in one relation that 'refers' to a tuple in another relation (like a 'logical pointer').

SQL Domain Constraints

- SQL supports various domain constraints to restrict attribute to valid domains
 - **NULL / NOT NULL** whether an attribute is allowed to become *NULL* (unknown)
 - **DEFAULT** to specify a default value
 - **CHECK(*condition*)** a Boolean *condition* that must hold for every tuple in the db instance

Example:

```
CREATE TABLE Student
(
    sid          INTEGER          PRIMARY KEY,
    name         VARCHAR(20)      NOT NULL,
    gender       CHAR              CHECK (gender IN ('M','F','T')),
    birthday     DATE              NULL,
    country      VARCHAR(20) ,
    level        INTEGER          DEFAULT 1 CHECK (level BETWEEN 1 and 5)
);
```

Data Loading / Storage

Data Storing

- Where are we now?
 - We have analysed our given data set
 - Cleaned it
 - Transformed it and created a corresponding relational database
- Next, we want to store the given data in our database.
- Main approaches:
 1. Command line tools
 2. Python loader
 3. (Combination of Python loader and stored procedures)

Approach 1: PSQL Data Loader

- Postgresql offers a command to load data directly from a CSV file into a database table

```
\COPY tablename FROM filename CSV [HEADER] [NULL '...']
```

- Many further options
 - Try `\help COPY`
- Pros:
 - Relatively fast and straight-forward
 - No programming needed
- Cons:
 - Only 1:1 mapping of CSV to tables; no data cleaning or transformation
 - Stops at the first error...

Approach 2: Python Loading Code

- Example: Creating a table and loading some data

```
# 1st: login to database|
conn = pgconnect()

# 2nd: ensure that the schema is in place
organisation_schema = """CREATE TABLE IF NOT EXISTS Organisation (
                           code VARCHAR(20) PRIMARY KEY,
                           orgName  VARCHAR(150)
                           )"""
pgexec (conn, organisation_schema, None, "Create Table Organisation")

# 3rd: Load data
# IMPORTANT: make sure the header line of CSV is without spaces!
insert_stmt = """INSERT INTO Organisation(code,orgName)
                 VALUES %(Code)s, %(Organisation)s)"""
for row in data_organisations:
    pgexec (conn, insert_stmt, row, "row inserted")
```

- Pros: Full flexibility; data cleaning and transformation possible
- Cons: Has to be hand-coded for each case

SQL DML Statements

- Insertion of new data into a table / relation
 - **Syntax:**
INSERT INTO table [“(”list-of-columns“)”] **VALUES** (“ list-of-expression “)”
 - Example:
INSERT INTO Students (sid, name) VALUES (53688, 'Smith')
- Updating of tuples in a table / relation
 - **Syntax:**
UPDATE table **SET** column“=“expression {“,”column“=“expression}
[**WHERE** search_condition]
 - Example: **UPDATE students**
SET gpa = gpa - 0.1
WHERE gpa >= 3.3
- Deleting of tuples from a table / relation
 - **Syntax:**
DELETE FROM table [**WHERE** search_condition]
 - Example:
DELETE FROM Students WHERE name = 'Smith'

Database Creation

– The full SQL schema for the given data model

```
DROP TABLE IF EXISTS Organisation CASCADE;  
CREATE TABLE IF NOT EXISTS Organisation (  
    code VARCHAR(20) PRIMARY KEY,  
    organisation VARCHAR(150)  
);
```

```
DROP TABLE IF EXISTS Station CASCADE;  
CREATE TABLE IF NOT EXISTS Station (  
    station VARCHAR(50) PRIMARY KEY,  
    siteName VARCHAR(50),  
    commence Date,  
    orgCode VARCHAR(50),  
    CONSTRAINT orgCodeFK  
        FOREIGN KEY (orgCode)  
        REFERENCES Organisation (code)  
);
```

```
DROP TABLE IF EXISTS Sensor CASCADE;  
CREATE TABLE IF NOT EXISTS Sensor (  
    sensor VARCHAR(20) PRIMARY KEY,  
    description VARCHAR(150) ,  
    metric VARCHAR(20)  
);
```

```
DROP TABLE IF EXISTS Measurement CASCADE;  
CREATE TABLE IF NOT EXISTS Measurement (  
    station VARCHAR(20),  
    sensor VARCHAR(20),  
    date DATE,  
    value FLOAT,  
    CONSTRAINT stationFK  
        FOREIGN KEY (station)  
        REFERENCES Station (Station),  
    CONSTRAINT sensorFK  
        FOREIGN KEY (sensor)  
        REFERENCES sensor (sensor)  
);
```


Review

Reprise Participation Marking

Requirements

- Submit code at end of each week
- **Jupyter Notebooks:**
 - The various exercises have placeholder cells marked as TODO:

```
# TODO: replace the content of this cell  
raise NotImplementedError
```

- The content of these cells needs to be replaced with your own solution
=> basis for participation marking

Output

- Code/spreadsheets from exercises

Marking

- 10% of overall mark
- each week's participation assessed as:
all done, partially done, no participation

Submit your Jupyter notebooks in Canvas by next Friday

Tips and Tricks

- Real data is 'dirty' – **data cleaning** and transformation essential
- Database systems are great for *shared, persistent* storage of *structured* data, and also for consistent updating ('life' data)
- But some caveats:
 - Schema-first
 - Relational model quite restrictive (1NF, no lists, collections etc)
 - Not too intuitive; 1:1 mapping from spreadsheets doomed to fail
 - Type-mismatches between programming languages and SQL
 - Needs to be installed and maintained
(though much better nowadays for SQLite and PostgreSQL)
- What's the benefit?
 - **Sharing!** Large data sets! Querying will give us some leverage too.

Next Time

Next Lecture : Querying and Summarising Data

Objective

To be able to extract a data set from a database, as well as to leverage on the SQL capabilities for in-database data summarisation and analysis.

Lecture

- Data Gathering reprise
- SQL querying
- Summarising data with SQL
- Statistic functions support in SQL

Readings

- Data Science from Scratch, Ch 23

Exercises

- [TODO]

TODO in W5

- Finish Grok Python modules
- Finish Grok SQL modules
- project data

Many Good Python Resources

- Hard to make recommendations given different backgrounds
- Look online, there are many free resources and example code
- A few lists:
 - <https://www.fullstackpython.com/best-python-resources.html>
 - <https://www.quora.com/Learning-Python/How-should-I-start-learning-Python-1>

Project Stage 1

Project Stage 1: Explore, Clean, Pitch

Objective

Explore a data set and define a research question based on research/business requirement.

Activities

- Choose a data set
- Explore and summarise data set
- Clean and prepare data
- Define problem

Output

- **2-page report** summarising data, problem and explorative analysis
 - how did you acquire the data?
 - which tools did you use to clean and explore the data set?
 - with title page & references: max 4p

Marking

- 13% of overall mark

SUGGESTED Timeline for Project Stage 1

- W1: Identify possible data sets
- W2: Identify and Explore possible data sets
- W3: Select project data set, define problem, complete exploration
- **W4: *Draft summary (problem & exploratory analysis)***
- W5: Clean and prepare data
- W6: Descriptive Stats, justification of suitability